

AMANAHA AI SOURCES & EVIDENCE REGISTER

Building Resilience Against AI-Generated Islamic Misinformation

Second Stage of competition source document | 5 October 2026

Purpose and scope

This register documents the external religious/evidence sources used by the AMANAHA AI MVP and clarifies how they are used. AMANAHA is an educational digital-resilience prototype: it checks source traceability, context and trust-risk signals, and provides controlled training scenarios. It does not issue fatwas, certify religious claims as universally true/false, or claim universal deepfake detection.

1. Primary source: Quranic verification principle

Quran 49:6 - Surah Al-Hujurat

Used as a general principle supporting verification before acting on or sharing information. In AMANAHA, this verse is not treated as proof that a particular user-submitted religious claim is correct or incorrect.

Reference: [Quran.com - Al-Hujurat 49:6](https://quran.com/49/6)

Role in MVP: Evidence/provenance card and verification-oriented learning feedback, including controlled scenarios where source verification or responsible sharing is the learning objective.

2. Institutional source: AI applications and fatwas

Egypt's Dar Al-Ifta - Using AI Applications to Obtain Fatwas

Published 2 December 2025. This source explains limitations of relying on current AI applications for fatwas and emphasises the scholarly knowledge, methodology and contextual understanding required for responsible fatwa practice.

Reference: [Dar Al-Ifta - Using AI Applications to Obtain Fatwas](#)

Role in MVP: Supports AMANAHA's safety boundary around AI-generated religious rulings and the training example in which a confident AI answer is presented without traceable evidence. It is used to encourage verification and qualified guidance where appropriate, not to automate a fatwa.

3. Institutional source: context and AI-generated fatwas

Egypt's Dar Al-Ifta - Resolving Family Issues: The Dilemma of AI-Generated Fatwas

Published 24 August 2026. The article distinguishes rapid AI-generated information from authoritative

religious guidance and discusses the importance of context, qualified expertise and responsible use of AI in sensitive religious and family matters.

Reference: [Dar Al-Ifta - The Dilemma of AI-Generated Fatwas](#)

Role in MVP: Supports context-awareness training and the scenario showing that even traceable religious material may not settle an individual case without considering circumstances and appropriate expertise.

4. Evidence-to-feature mapping

MVP area	Source/evidence	How it is used	Boundary
AMANAHA CHECK	Quran 49:6; Dar Al-Ifta AI/fatwa guidance	Prompts verification, displays source traceability/evidence status, and explains why unsupported religious claims require caution.	No automated halal/haram or true/false religious verdict.
TRAIN - AI Over-Reliance	Dar Al-Ifta: Using AI Applications to Obtain Fatwas	Supports feedback that confident AI output is not itself evidence of religious authority.	Controlled educational scenario; not a live fatwa evaluation.
TRAIN - Source Verification / Responsible Sharing	Quran 49:6	Supports the general learning principle of verifying information before relying on or sharing it.	Verse is not used as proof of the truth/falsity of a specific claim.
TRAIN - Context Awareness	Dar Al-Ifta: Dilemma of AI-Generated Fatwas	Supports the lesson that source authenticity alone may not resolve a sensitive individual case without context.	No personalised religious ruling is generated.
TRAIN - Manipulated Content	AMANAHA controlled verification protocol	Teaches users to seek provenance/original upload/full context before treating a convincing clip as authentic.	This is an internal training method, not external religious evidence and not a deepfake detector.

5. Controlled and simulated content used in the MVP

The five training scenarios are controlled/simulated learning examples created for the AMANAHA MVP. Fictional handles, simulated quote cards, short-clip previews and forwarded-audio examples are not presented as real people, real social-media posts, real engagement statistics or verified incidents.

The five learning dimensions are: Source Verification; AI Over-Reliance; Context Awareness; Manipulated Content Awareness; Responsible Sharing.

The session profile and scores are learning indicators generated from the user's decisions during that session. They are not validated psychometric measures and are not presented as clinical, scientific or population-level scores.

6. Technical disclosure relevant to sources

- The competition MVP uses a browser-based HTML/CSS/JavaScript frontend hosted on GitHub Pages.
- AMANAHA CHECK now uses a live, controlled retrieval-augmented generation (RAG) pipeline. Relevant evidence is selected from an approved source collection before generation, and only the retrieved evidence is supplied to the language model for the trust-risk explanation.
- The generative layer runs through a Cloudflare Worker using Cloudflare Workers AI. The current deployed model is @cf/meta/llama-3.1-8b-instruct-fast.
- The RAG layer is intentionally constrained: it explains source traceability, evidence sufficiency and trust-risk signals; it does not issue fatwas, decide halal/haram, or certify a religious claim as definitively true or false.
- Where no sufficiently relevant approved evidence is retrieved, AMANAHA abstains from generating a religious judgement. If the AI service or structured output fails, the application falls back to deterministic local safety checks and continues to display retrieved evidence rather than inventing an answer.
- The MVP has no user database or user authentication. No AI-provider secret or API key is exposed in the public frontend repository; the AI binding is handled server-side by the Cloudflare Worker.
- External source links are provided for provenance where applicable; source availability remains subject to the originating websites.
- The current scenario set has not been independently scholar-reviewed or psychometrically validated. Expert content validation and user evaluation remain future validation steps.

7. Source-selection and reliability principles

For this MVP, religious/evidence references are limited to directly traceable primary or institutional sources used for the specific learning point. AMANAHA separates source provenance from religious adjudication: a traceable source may still require context, and absence of a traceable source does not by itself establish that a claim is false. Where evidence is insufficient, the product is designed to say so and encourage checking the original source or appropriate qualified guidance.

Technical implementation note: Cloudflare Workers and Workers AI are used as runtime infrastructure for the server-side grounded generation layer. They are implementation technologies, not religious evidence sources.

8. Reference list

1. Quran.com. Surah Al-Hujurat 49:6. <https://quran.com/en/49/6> Accessed 5 October 2026.
2. Egypt's Dar Al-Ifta. Using AI Applications to Obtain Fatwas. 2 December 2025. <https://www.dar-alifta.org/en/fatwa/details/22255/using-ai-applications-to-obtain-fatwas> Accessed 5 October 2026.
3. Egypt's Dar Al-Ifta. Resolving Family Issues: The Dilemma of AI-Generated Fatwas. 24 August 2026. <https://www.dar-alifta.org/en/article/details/12325/resolving-family-issues-the-dilemma-of-ai-generated-fatwas> Accessed 5 October 2026.

Important interpretation note: The references above support AMANAHA's verification, context and responsible-use learning principles. They should not be interpreted as endorsements of AMANAHA AI by Quran.com, Egypt's Dar Al-Ifta, or any other institution.